

CHAPTER 43

RED-BLACK TREES

Objectives

- To know what a red-black tree is (§43.1).
- To convert a red-black tree to a 2–4 tree and vice versa (§43.2).
- To design the **RBTree** class that extends the **BST** class (§43.3).
- To insert an element in a red-black tree and resolve the double-red violation if necessary (§43.4).
- To delete an element from a red-black tree and resolve the double-black problem if necessary (§43.5).
- To implement and test the **RBTree** class (§§43.6–43.7).
- To compare the performance of AVL trees, 2–4 trees, and **RBTree** (§43.8).





43.1 Introduction

A red-black tree is a balanced binary search tree derived from a 2-4 tree. A red-black tree corresponds to a 2-4 tree.

Each node in a red-black tree has a *color attribute* red or black, as shown in Figure 43.1(a). A node is called *external* if its left or right subtree is empty. Note that a leaf node is external, but an external node is not necessarily a leaf node. For example, node 25 is external, but it is not a leaf. The *black depth* of a node is defined as the number of black nodes in a path from the node to the root. For example, the black depth of node 25 is 2, and that of node 27 is 2.



Note

The red nodes appear in blue in the text.

A red-black tree has the following properties:

1. The root is black.
2. Two adjacent nodes cannot be both red.
3. All external nodes have the same black depth.

The red-black tree in Figure 43.1(a) satisfies all three properties. A red-black tree can be converted to a 2-4 tree, and vice versa. Figure 43.1(b) shows an equivalent 2-4 tree for the red-black tree in Figure 43.1(a).




FIGURE 43.1 A red-black tree can be represented using a 2-4 tree, and vice versa.

43.2 Conversion between Red-Black Trees and 2-4 Trees



This section discusses the correspondence between a red-black tree and a 2-4 tree.

You can design insertion and deletion algorithms for red-black trees without having knowledge of 2-4 trees. However, the correspondence between red-black trees and 2-4 trees provides useful intuition about the structure of red-black trees and operations. For this reason, this section discusses the correspondence between these two types of trees.

To convert a red-black tree to a 2-4 tree, simply merge every red node with its parent to create a 3-node or a 4-node. For example, the red nodes 15 and 34 are merged to their parent to create a 4-node, and the red node 27 is merged to its parent to create a 3-node, as shown in Figure 43.1(b).

To convert a 2-4 tree to a red-black tree, perform the following transformations for each node u :

1. If u is a 2-node, color it black, as shown in Figure 43.2(a).
2. If u is a 3-node with element values e_0 and e_1 , there are two ways to convert it. Either make e_0 the parent of e_1 or make e_1 the parent of e_0 . In any case, color the parent black and the child red, as shown in Figure 43.2(b).
3. If u is a 4-node with element values e_0 , e_1 , and e_2 , make e_1 the parent of e_0 and e_2 . Color e_1 black and e_0 and e_2 red, as shown in Figure 43.2(c).




FIGURE 43.2 A node in a 2-4 tree can be transformed to nodes in a red-black tree.

Let us apply the transformation for the 2-4 tree in Figure 43.1(b). After transforming the 4-node, the tree is as shown in Figure 43.3(a). After transforming the 3-node, the tree is as shown in Figure 43.3(b). Note the transformation for a 3-node is not unique. Therefore, the conversion from a 2-4 tree to a red-black tree is *not unique*. After transforming the 3-node, the tree could also be as shown in Figure 43.3(c).




FIGURE 43.3 The conversion from a 2-4 tree to a red-black tree is not unique.

You can prove the conversion results in a red-black tree that satisfies all three properties.

Property 1. The root is black.

Proof: If the root of a 2-4 tree is a 2-node, the root of the red-black tree is black. If the root of a 2-4 tree is a 3-node or 4-node, the transformation produces a black parent at the root.

Property 2. Two adjacent nodes cannot be both red.

Proof: Since the parent of a red node is always black, no two adjacent nodes can be both red.

Property 3. All external nodes have the same black depth.

Proof: When you convert a node in a 2-4 tree to red-black tree nodes, you get one black node and zero, one, or two red nodes as its children, depending on whether the original node is a 2-, 3-, or 4-node. Only a leaf 2-4 node may produce external red-black nodes. Since a 2-4 tree is perfectly balanced, the number of black nodes in any path from the root to an external node is the same.



- 43.2.1 What is a red-black tree? What is an external node? What is black depth?
- 43.2.2 Describe the properties of a red-black tree.
- 43.2.3 How do you convert a red-black tree to a 2-4 tree? Is the conversion unique?
- 43.2.4 How do you convert a 2-4 tree to a red-black tree? Is the conversion unique?



43.3 Designing Classes for Red-Black Trees

A red-black tree designs a class for a red-black tree.

A red-black tree is a binary search tree. So, you can define the **RBTree** class to extend the **BST** class, as shown in Figure 43.4. The **BST** and **TreeNode** classes are defined in §26.2.5.

Each node in a red-black tree has a color property. Because the color is either red or black, it is efficient to use the **boolean** type to denote it. The **RBTreeNode** class can be defined to extend **BST.TreeNode** with the color property. For convenience, we also provide the methods for checking the color and setting a new color. Note that **TreeNode** is defined as a static inner class in **BST**. **RBTreeNode** will be defined as a static inner class in **RBTree**. Note that **BSTNode** contains the data fields **element**, **left**, and **right**, which are inherited in **RBTreeNode**. So, **RBTreeNode** contains four data fields, as pictured in Figure 43.5.

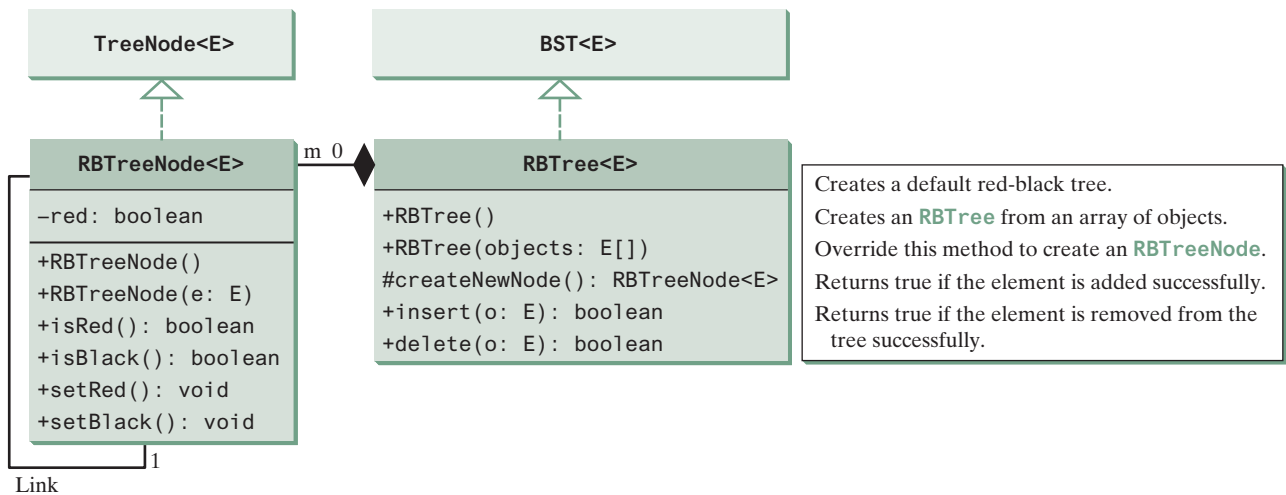


FIGURE 43.4 The **RBTree** class extends **BST** with new implementations for the **insert** and **delete** methods.

node: <code>RBTreeNode<E></code>
<code>#element: E</code>
<code>-red: boolean</code>
<code>#left: TreeNode</code>
<code>#right: TreeNode</code>

FIGURE 43.5 An `RBTreeNode` contains data fields `element`, `red`, `left`, and `right`.

In the `BST` class, the `createNewNode()` method creates a `TreeNode` object. This method is overridden in the `RBTree` class to create an `RBTreeNode`. Note the return type of the `createNewNode()` method in the `BST` class is `TreeNode`, but the return type of the `createNewNode()` method in `RBTree` class is `RBTreeNode`. This is fine, since `RBTreeNode` is a subtype of `TreeNode`.

Searching an element in a red-black tree is the same as searching in a regular binary search tree. So, the `search` method defined in the `BST` class also works for `RBTree`.

The `insert` and `delete` methods are overridden to insert and delete an element and perform operations for coloring and restructuring if necessary to ensure that the three properties of the red-black tree are satisfied.



Pedagogical Note

Run from <http://liveexample.pearsoncmg.com/dsanimation/RBTree.html> to see how a red-black tree works, as shown in Figure 43.6.

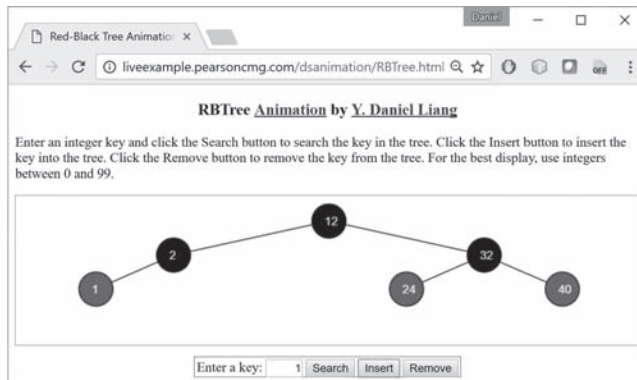


FIGURE 43.6 The animation tool enables you to insert, delete, and search elements in a red-black tree visually.

43.4 Overriding the `insert` Method

This section discusses how to insert an element to red-black tree.

A new element is always inserted as a leaf node. If the new node is the root, color it black. Otherwise, color it red. If the parent of the new node is red, it violates Property 2 of the red-black tree. We call this a *double-red* violation.

Let u denote the new node inserted, v the parent of u , w the parent of v , and x the sibling of v . To fix the double-red violation, consider two cases:

Case 1: x is black or x is null. There are four possible configurations for u , v , w , and x , as shown in Figures 43.7(a), 43.8(a), 43.9(a), and 43.10(a). In this case, u , v , and w form a 4-node in the corresponding 2-4 tree, as shown in Figures 43.7(c), 43.8(c), 43.9(c), and 43.10(c), but are represented incorrectly in the red-black tree. To correct this error, restructure and recolor three nodes u , v , and w , as shown in Figures 43.7(b), 43.8(b), 43.9(b), and 43.10(b). Note x , y_1 , y_2 , and y_3 may be null.



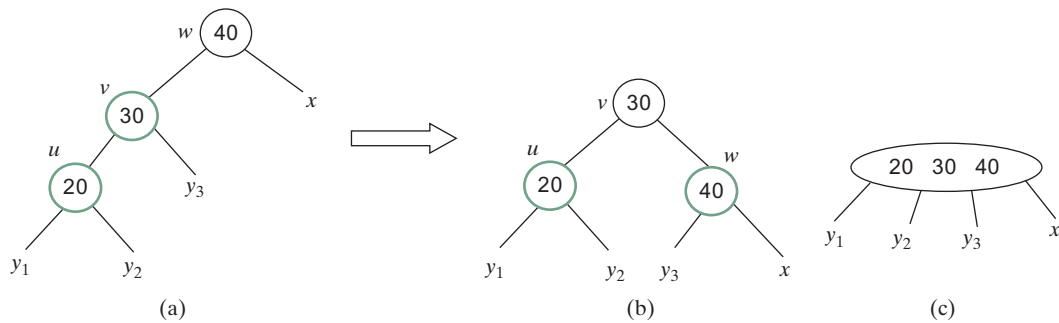


FIGURE 43.7 Case 1.1: $u < v < w$.

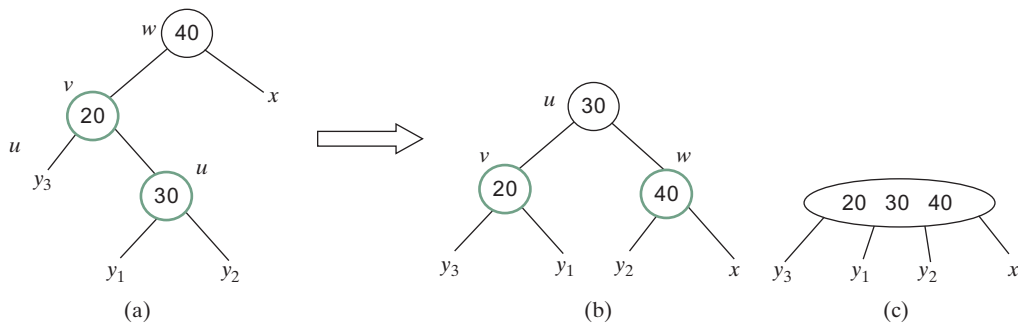


FIGURE 43.8 Case 1.2: $v < u < w$




FIGURE 43.9 Case 1.3: $w < v < u$




FIGURE 43.10 Case 1.4: $w < u < v$

Case 2: x is red. There are four possible configurations for u , v , w , w , and x , as shown in Figures 43.11(a), 43.11(b), 43.11(c), and 43.11(d). All of these configurations correspond to an overflow situation in the corresponding 4-node in a 2-4 tree, as shown in Figure 43.12(a). A splitting operation is performed to fix the overflow problem in a 2-4 tree, as shown in Figure 43.12(b). We perform an equivalent recoloring operation to fix the problem in a red-black tree. Color w and u red and color two children of w black. Assume u is a left child of v , as shown




FIGURE 43.11 Case 2 has four possible configurations.

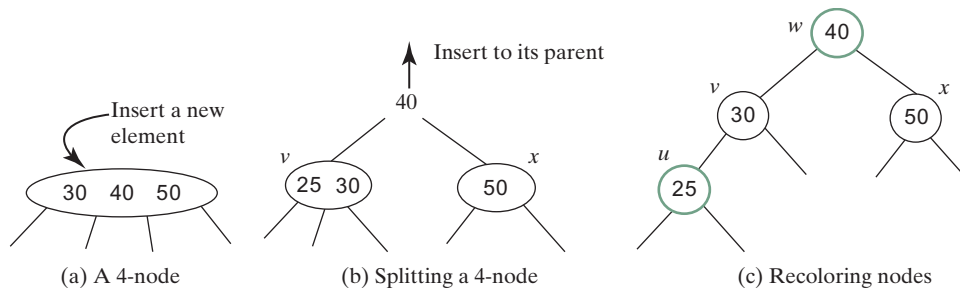


FIGURE 43.12 Splitting a 4-node corresponds to recoloring the nodes in the red-black tree.

in Figure 43.11(a). After recoloring, the nodes are shown in Figure 43.12(c). Now w is red, if w 's parent is black, the double-red violation is fixed. Otherwise, a new double-red violation occurs at node w . We need to continue the same process to eliminate the double-red violation at w , recursively.

A more detailed algorithm for inserting an element is described in Listing 43.1.

LISTING 43.1 Inserting an Element to a Red-Black Tree

```

1 public boolean insert(E e) {
2     boolean successful = super.insert(e);
3     if (!successful)
4         return false; // e is already in the tree
5     else {
6         ensureRBTree(e);
7     }
8
9     return true; // e is inserted
10 }
11
12 /** Ensure that the tree is a red-black tree */
13 private void ensureRBTree(E e) {
14     Get the path that leads to element e from the root.

```

```

15     int i = path.size() - 1; // Index to the current node in the path
16     Get u, v from the path. u is the node that contains e and v
17         is the parent of u.
18     Color u red;
19
20     if (u == root) // If e is inserted as the root, set root black
21         u.setBlack();
22     else if (v.isRed())
23         fixDoubleRed(u, v, path, i); // Fix double-red violation at u
24 }
25
26 /** Fix double-red violation at node u */
27 private void fixDoubleRed(RBTreeNode<E> u, RBTreeNode<E> v,
28     ArrayList<TreeNode<E>> path, int i) {
29     Get w from the path. w is the grandparent of u.
30
31     // Get v's sibling named x
32     RBTreeNode<E> x = (w.left == v) ?
33         (RBTreeNode<E>)(w.right) : (RBTreeNode<E>)(w.left);
34
35     if (x == null || x.isBlack()) {
36         // Case 1: v's sibling x is black
37         if (w.left == v && v.left == u) {
38             // Case 1.1: u < v < w, Restructure and recolor nodes
39         }
40         else if (w.left == v && v.right == u) {
41             // Case 1.2: v < u < w, Restructure and recolor nodes
42         }
43         else if (w.right == v && v.right == u) {
44             // Case 1.3: w < v < u, Restructure and recolor nodes
45         }
46         else {
47             // Case 1.4: w < u < v, Restructure and recolor nodes
48         }
49     }
50     else { // Case 2: v's sibling x is red
51         Color w and u red
52         Color two children of w black.
53
54         if (w is root) {
55             Set w black;
56         }
57         else if (the parent of w is red) {
58             // Propagate along the path to fix new double-red violation
59             u = w;
60             v = parent of w;
61             fixDoubleRed(u, v, path, i - 2); // i - 2 propagates upward
62         }
63     }
64 }

```

The **insert(E e)** method (lines 1–10) invokes the **insert** method in the **BST** class to create a new leaf node for the element (line 2). If the element is already in the tree, return false (line 4). Otherwise, invoke **ensureRBTree(e)** (line 6) to ensure that the tree satisfies the color and black depth property of the red-black tree.

The **ensureRBTree(E e)** method (lines 13–24) obtains the path that leads to **e** from the root (line 14), as shown in Figure 43.13. This path plays an important role to implement the algorithm. From this path, you get nodes **u** and **v** (lines 16–17). If **u** is the root, color **u** black (lines 20–21). If **v** is red, a double-red violation occurs at node **u**. Invoke **fixDoubleRed** to fix the problem.




FIGURE 43.13 The path consists of the nodes from u to the root.

The `fixDoubleRed` method (lines 27–63) fixes the double-red violation. It first obtains w (the parent of v) from the path (line 29) and x (the sibling of v) (lines 32–33). If x is empty or a black node, restructure and recolor three nodes u , v , and w to eliminate the problem (lines 35–49). If x is a red node, recolor the nodes u , v , w , and x (lines 51–52). If w is the root, color w black (lines 54–56). If the parent of w is red, the double-red violation reappears at w . Invoke `fixDoubleRed` with new u and v to fix the problem (line 61). Note that now $i - 2$ points to the new u in the path. This adjustment is necessary to locate the new nodes w and parent of w along the path.

Figure 43.14 shows the steps of inserting 34, 3, 50, 20, 15, 16, 25, and 27 into an empty red-black tree. When inserting 20 into the tree in (d), Case 2 applies to recolor 3 and 50 to black. When inserting 15 into the tree in (g), Case 1.4 applies to restructure and recolor nodes 15, 20, and 3. When inserting 16 into the tree in (i), Case 2 applies to recolor nodes 3 and 20 to black and nodes 15 and 16 to red. When inserting 27 into the tree in (l), Case 2 applies to recolor nodes 16 and 25 to black and nodes 20 and 27 to red. Now a new double-red problem occurs at node 20. Apply Case 1.2 to restructure and recolor nodes. The new tree is shown in (n).

43.5 Overriding the `delete` Method

This section discusses how to delete an element to red-black tree.

To delete an element from a red-black tree, first search the element in the tree to locate the node that contains the element. If the element is not in the tree, the method returns false. Let u be the node that contains the element. If u is an internal node with both left and right children, find the rightmost node in the left subtree of u . Replace the element in u with the element in the rightmost node. Now we will only consider deleting external nodes.

Let u be an external node to be deleted. Since u is an external node, it has at most one child, denoted by `childOfu`. `childOfu` may be `null`. Let `parentOfu` denote the parent of u , as shown in Figure 43.15(a). Delete u by connecting `childOfu` with `parentOfu`, as shown in Figure 43.15(b).

Consider the following case:

- If u is red, we are done.
- If u is black and `childOfu` is red, color `childOfu` black to maintain the black height for `childOfu`.
- Otherwise, assign `childOfu` a fictitious *double black*, as shown in Figure 43.16(a). We call this a *double-black problem*, which indicates that the black depth is short by 1, caused by deleting a black node u .






FIGURE 43.14 Inserting into a red-black tree: (a) initial empty tree; (b) inserting **34**; (c) inserting **3**; (d) inserting **50**; (e) inserting **20** causes a double red; (f) after recoloring (Case 2); (g) inserting **15** causes a double red; (h) after restructuring and recoloring (Case 1.4); (i) inserting **16** causes a double red; (j) after recoloring (Case 2); (k) inserting **25**; (l) inserting **27** causes a double red at **27**; (m) a double red at 20 reappears after recoloring (Case 2); and (n) after restructuring and recoloring (Case 1.2).




FIGURE 43.15 u is an external node and `childOfu` may be null.




FIGURE 43.16 (a) `childOfu` is denoted double black. (b) u corresponds to an empty node in a 2-4 tree.

A double black in a red-black tree corresponds to an empty node for u (i.e., underflow situation) in the corresponding 2-4 tree, as shown in Figure 43.16(b). To fix the double-black problem, we will perform equivalent transfer and fusion operations. Consider three cases:

Case 1: The sibling y of `childOfu` is black and has a red child. This case has four possible configurations, as shown in Figures 43.17(a), 43.18(a), 43.19(a), and 43.20(a). The dashed circle denotes that the node is either red or black. To eliminate the double-black problem, restructure and recolor the nodes, as shown in Figures 43.17(b), 43.18(b), 43.19(b), and 43.20(b).




FIGURE 43.17 Case 1.1: The sibling y of `childOfu` is black and y_1 is red.




FIGURE 43.18 Case 1.2: The sibling y of `childOfu` is black and y_2 is red.




FIGURE 43.19 Case 1.3: The sibling *y* of *childOfu* is black and *y1* is red.




FIGURE 43.20 Case 1.4: the sibling *y* of *childOfu* is black and *y2* is red.



Note

Case I corresponds to a *transfer* operation in the 2-4 tree. For example, the corresponding 2-4 tree for Figure 43.17(a) is shown in Figure 43.21(a), and it is transformed into Figure 43.21(b) through a transfer operation.




FIGURE 43.21 Case I corresponds to a transfer operation in the corresponding 2-4 tree.

Case 2: The sibling *y* of *childOfu* is black and its children are black or *null*. In this case, change *y*'s color to red. If *parent* is red, change it to black, and we are done, as shown in Figure 43.22. If *parent* is black, we denote *parent* double black, as shown in Figure 43.23. The double-black problem *propagates* to the parent node.




FIGURE 43.22 Case 2: Recoloring eliminates the double-black problem if *parent* is red.




FIGURE 43.23 Case 2: Recoloring propagates the double-black problem if `parent` is black.



Note

Figures 43.22 and 43.23 show that `childOfu` is a right child of `parent`. If `childOfu` is a left child of `parent`, recoloring is performed identically.



Note

Case 2 corresponds to a *fusion* operation in the 2-4 tree. For example, the corresponding 2-4 tree for Figure 43.22(a) is shown in Figure 43.24(a), and it is transformed into Figure 43.24(b) through a fusion operation.

Case 3: The sibling `y` of `childOfu` is red. In this case, perform an *adjustment* operation. If `y` is a left child of `parent`, let `y1` and `y2` be the left and right children of `y`, as shown in Figure 43.25. If `y` is a right child of `parent`, let `y1` and `y2` be the left and right child of `y`, as shown in Figure 43.26. In both cases, color `y` black and `parent` red. `childOfu` is still a fictitious double-black node. After the adjustment, the sibling of `childOfu` is now black, and either Case 1 or Case 2 applies. If Case 1 applies, a one-time restructuring and recoloring operation eliminates the double-black problem. If Case 2 applies, the double-black problem cannot reappear, since `parent` is now red. Therefore, one-time application of Case 1 or Case 2 will complete Case 3.



Note

Case 3 results from the fact that a 3-node may be transformed in two ways to a red-black tree, as shown in Figure 43.27.




FIGURE 43.24 Case 2 corresponds to a fusion operation in the corresponding 2-4 tree.




FIGURE 43.25 Case 3.1: `y` is a left red child of `parent`.




FIGURE 43.26 Case 3.2: y is a right red child of $parent$.




FIGURE 43.27 A 3-node may be transformed in two ways to red-black tree nodes.

Based on the foregoing discussion, Listing 43.2 presents a more detailed algorithm for deleting an element.

LISTING 43.2 Deleting an Element from a Red-Black Tree

```

1  public boolean delete(E e) {
2      Locate the node to be deleted
3      if (the node is not found)
4          return false;
5
6      if (the node is an internal node) {
7          Find the rightmost node in the subtree of the node;
8          Replace the element in the node with the one in rightmost;
9          The rightmost node is the node to be deleted now;
10     }
11
12     Obtain the path from the root to the node to be deleted;
13
14     // Delete the last node in the path and propagate if needed
15     deleteLastNodeInPath(path);
16 
```

```

17     size--; // After one element deleted
18     return true; // Element deleted
19 }
20
21 /** Delete the last node from the path. */
22 public void deleteLastNodeInPath(ArrayList<TreeNode<E>> path) {
23     Get the last node u in the path;
24     Get parentOfu and grandparentOfu in the path;
25     Get childOfu from u;
26     Delete node u. Connect childOfu with parentOfu
27
28     // Recolor the nodes and fix double black if needed
29     if (childOfu == root || u.isRed())
30         return; // Done if childOfu is root or if u is red
31     else if (childOfu != null && childOfu.isRed())
32         childOfu.setBlack(); // Set it black, done
33     else // u is black, childOfu is null or black
34         // Fix double black on parentOfu
35         fixDoubleBlack(grandparentOfu, parentOfu, childOfu, path, i);
36 }
37
38 /** Fix the double black problem at node parent */
39 private void fixDoubleBlack(
40     RBTreeNode<E> grandparent, RBTreeNode<E> parent,
41     RBTreeNode<E> db, ArrayList<TreeNode<E>> path, int i) {
42     Obtain y, y1, and y2
43
44     if (y.isBlack() && y1 != null && y1.isRed()) {
45         if (parent.right == db) {
46             // Case 1.1: y is a left black sibling and y1 is red
47             Restructure and recolor parent, y, and y1 to fix the problem;
48         }
49         else {
50             // Case 1.3: y is a right black sibling and y1 is red
51             Restructure and recolor parent, y1, and y to fix the problem;
52         }
53     }
54     else if (y.isBlack() && y2 != null && y2.isRed()) {
55         if (parent.right == db) {
56             // Case 1.2: y is a left black sibling and y2 is red
57             Restructure and recolor parent, y2, and y to fix the problem;
58         }
59         else {
60             // Case 1.4: y is a right black sibling and y2 is red
61             Restructure and recolor parent, y, and y2 to fix the problem;
62         }
63     }
64     else if (y.isBlack()) {
65         // Case 2: y is black and y's children are black or null
66         Recolor y to red;
67
68         if (parent.isRed())
69             parent.setBlack(); // Done
70         else if (parent != root) {
71             // Propagate double black to the parent node
72             // Fix new appearance of double black recursively
73             db = parent;
74             parent = grandparent;
75             grandparent =
76                 (i >= 3) ? (RBTreeNode<E>)(path.get(i - 3)) : null;

```

```

77         fixDoubleBlack(grandparent, parent, db, path, i - 1);
78     }
79 }
80 else if (y.isRed()) {
81     if (parent.right == db) {
82         // Case 3.1: y is a left red child of parent
83         parent.left = y2;
84         y.right = parent;
85     }
86     else {
87         // Case 3.2: y is a right red child of parent
88         parent.right = y.left;
89         y.left = parent;
90     }
91     parent.setRed(); // Color parent red
92     y.setBlack(); // Color y black
93     connectNewParent(grandparent, parent, y); // y is new parent
94     fixDoubleBlack(y, parent, db, path, i - 1);
95 }
96 }
97 }

```

The `delete(E e)` method (lines 1–19) locates the node that contains `e` (line 2). If the node does not exist, return `false` (lines 3–4). If the node is an internal node, find the right most node in its left subtree and replace the element in the node with the element in the right most node (lines 6–9). Now the node to be deleted is an external node. Obtain the path from the root to the node (line 12). Invoke `deleteLastNodeInPath(path)` to delete the last node in the path and ensure that the tree is still a red-black tree (line 15).

The `deleteLastNodeInPath` method (lines 22–36) obtains the last node `u`, `parentOfu`, `grandparentOfu`, and `childOfu` (lines 23–26). If `childOfu` is the root or `u` is red, the tree is fine (lines 29–30). If `childOfu` is red, color it black (lines 31–32). We are done. Otherwise, `u` is black and `childOfu` is `null` or black. Invoke `fixDoubleBlack` to eliminate the double-black problem (line 35).

The `fixDoubleBlack` method (lines 39–97) eliminates the double-black problem. Obtain `y`, `y1`, and `y2` (line 42). `y` is the sibling of the double-black node. `y1` and `y2` are the left and right children of `y`. Consider three cases:

1. If `y` is black and one of its children is red, the double-black problem can be fixed by one-time restructuring and recoloring in Case 1 (lines 44–63).
2. If `y` is black and its children are `null` or black, change `y` to red. If `parent` of `y` is black, denote `parent` to be the new double-black node and invoke `fixDoubleBlack` recursively (line 77).
3. If `y` is red, adjust the nodes to make `parent` a child of `y` (lines 84, 89) and color `parent` red and `y` black (lines 92–93). Make `y` the new parent (line 94). Recursively invoke `fixDoubleBlack` on the same double-black node with a different color for `parent` (line 95).

Figure 43.28 shows the steps of deleting elements. To delete `50` from the tree in Figure 43.28(a), apply Case 1.2, as shown in Figure 43.28(b). After restructuring and recoloring, the new tree is as shown in Figure 43.28(c).

When deleting `20` in Figure 43.28(c), `20` is an internal node, and it is replaced by `16`, as shown in Figure 43.28(d). Now Case 2 applies to deleting the rightmost node, as shown in Figure 43.28(e). Recolor the nodes results in a new tree, as shown in Figure 43.28(f).

When deleting `15`, connect node 3 with node 20 and color node 3 black, as shown in Figure 43.28(g). We are done.




FIGURE 43.28 Delete elements from a red-black tree.

After deleting **25**, the new tree is as shown in Figure 43.28(j). Now delete **16**. Apply Case 2, as shown in Figure 43.28(k). The new tree is shown in Figure 43.28(l).

After deleting **34**, the new tree is as shown in Figure 43.28(m).

After deleting **27**, the new tree is as shown in Figure 43.28(n).



43.5.1 What are the data fields in `RBTreeNode`?

43.5.2 How do you insert an element into a red-black tree and how do you fix the double-red violation?

43.5.3 How do you delete an element from a red-black tree and how do you fix the double-black problem?

43.5.4 Show the change of the tree when inserting **1, 2, 3, 4, 10, 9, 7, 5, 8**, and **6** into it, in this order.

43.5.5 For the tree built in the preceding question, show the change of the tree after deleting **1, 2, 3, 4, 10, 9, 7, 5, 8**, and **6** from it in this order.



43.6 Implementing `RBTree` Class

This section implements the `RBTree` class.

Listing 43.3 gives a complete implementation for the `RBTree` class.

LISTING 43.3 `RBTree.java`

```

1  import java.util.ArrayList;
2
3  public class RBTree<E extends Comparable<E>> extends BST<E> {
4      /** Create a default RB tree */
5      public RBTree() {
6      }
7
8      /** Create an RB tree from an array of elements */
9      public RBTree(E[] elements) {
10         super(elements);
11     }
12
13     @Override /** Override createNewNode to create an RBTreeNode */
14     protected RBTreeNode<E> createNewNode(E e) {
15         return new RBTreeNode<E>(e);
16     }
17
18     @Override /** Override the insert method to
19         balance the tree if necessary */
20     public boolean insert(E e) {
21         boolean successful = super.insert(e);
22         if (!successful)
23             return false; // e is already in the tree
24         else {
25             ensureRBTree(e);
26         }
27
28         return true; // e is inserted
29     }
30
31     /** Ensure that the tree is a red-black tree */
32     private void ensureRBTree(E e) {
33         // Get the path that leads to element e from the root
34         ArrayList<TreeNode<E>> path = path(e);
35
36         int i = path.size() - 1; // Index to the current node in the path
37
38         // u is the last node in the path. u contains element e
39         RBTreeNode<E> u = (RBTreeNode<E>)(path.get(i));

```

```

40
41 // v is the parent of of u, if exists
42 RBTTreeNode<E> v = (u == root) ? null :
43     (RBTTreeNode<E>)(path.get(i - 1));
44
45 u.setRed(); // It is OK to set u red
46
47 if (u == root) // If e is inserted as the root, set root black
48     u.setBlack();
49 else if (v.isRed())
50     fixDoubleRed(u, v, path, i); // Fix double-red violation at u
51 }
52
53 /** Fix double-red violation at node u */
54 private void fixDoubleRed(RBTTreeNode<E> u, RBTTreeNode<E> v,
55     ArrayList<TreeNode<E>> path, int i) {
56     // w is the grandparent of u
57     RBTTreeNode<E> w = (RBTTreeNode<E>)(path.get(i - 2));
58     RBTTreeNode<E> parentOfw = (w == root) ? null :
59         (RBTTreeNode<E>)path.get(i - 3);
60
61     // Get v's sibling named x
62     RBTTreeNode<E> x = (w.left == v) ?
63         (RBTTreeNode<E>)(w.right) : (RBTTreeNode<E>)(w.left);
64
65     if (x == null || x.isBlack()) {
66         // Case 1: v's sibling x is black
67         if (w.left == v && v.left == u) {
68             // Case 1.1: u < v < w, Restructure and recolor nodes
69             restructureRecolor(u, v, w, w, parentOfw);
70
71             w.left = v.right; // v.right is y3 in Figure 43.6
72             v.right = w;
73         }
74         else if (w.left == v && v.right == u) {
75             // Case 1.2: v < u < w, Restructure and recolor nodes
76             restructureRecolor(v, u, w, w, parentOfw);
77             v.right = u.left;
78             w.left = u.right;
79             u.left = v;
80             u.right = w;
81         }
82         else if (w.right == v && v.right == u) {
83             // Case 1.3: w < v < u, Restructure and recolor nodes
84             restructureRecolor(w, v, u, w, parentOfw);
85             w.right = v.left;
86             v.left = w;
87         }
88         else {
89             // Case 1.4: w < u < v, Restructure and recolor nodes
90             restructureRecolor(w, u, v, w, parentOfw);
91             w.right = u.left;
92             v.left = u.right;
93             u.left = w;
94             u.right = v;
95         }
96     }
97     else { // Case 2: v's sibling x is red
98         // Recolor nodes
99         w.setRed();
100        u.setRed();

```

```

101         ((RBTreeNode<E>)(w.left)).setBlack();
102         ((RBTreeNode<E>)(w.right)).setBlack();
103
104         if (w == root) {
105             w.setBlack();
106         }
107         else if (((RBTreeNode<E>)parentOfw).isRed()) {
108             // Propagate along the path to fix new double-red violation
109             u = w;
110             v = (RBTreeNode<E>)parentOfw;
111             fixDoubleRed(u, v, path, i - 2); // i - 2 propagates upward
112         }
113     }
114 }
115
116 /** Connect b with parentOfw and recolor a, b, c for a < b < c */
117 private void restructureRecolor(RBTreeNode<E> a, RBTreeNode<E> b,
118     RBTreeNode<E> c, RBTreeNode<E> w, RBTreeNode<E> parentOfw) {
119     if (parentOfw == null)
120         root = b;
121     else if (parentOfw.left == w)
122         parentOfw.left = b;
123     else
124         parentOfw.right = b;
125
126     b.setBlack(); // b becomes the root in the subtree
127     a.setRed(); // a becomes the left child of b
128     c.setRed(); // c becomes the right child of b
129 }
130
131 @Override /** Delete an element from the RBTree.
132     * Return true if the element is deleted successfully
133     * Return false if the element is not in the tree */
134 public boolean delete(E e) {
135     // Locate the node to be deleted
136     TreeNode<E> current = root;
137     while (current != null) {
138         if (e.compareTo(current.element) < 0) {
139             current = current.left;
140         }
141         else if (e.compareTo(current.element) > 0) {
142             current = current.right;
143         }
144         else
145             break; // Element is in the tree pointed by current
146     }
147
148     if (current == null)
149         return false; // Element is not in the tree
150
151     java.util.ArrayList<TreeNode<E>> path;
152
153     // current node is an internal node
154     if (current.left != null && current.right != null) {
155         // Locate the rightmost node in the left subtree of current
156         TreeNode<E> rightMost = current.left;
157         while (rightMost.right != null) {
158             rightMost = rightMost.right; // Keep going to the right
159         }
160

```

```

161     path = path(rightMost.element); // Get path before replacement
162
163     // Replace the element in current by the element in rightMost
164     current.element = rightMost.element;
165 }
166 else
167     path = path(e); // Get path to current node
168
169 // Delete the last node in the path and propagate if needed
170 deleteLastNodeInPath(path);
171
172 size--; // After one element deleted
173 return true; // Element deleted
174 }
175
176 /** Delete the last node from the path. */
177 public void deleteLastNodeInPath(ArrayList<TreeNode<E>> path) {
178     int i = path.size() - 1; // Index to the node in the path
179     // u is the last node in the path
180     RBTreeNode<E> u = (RBTreeNode<E>)(path.get(i));
181     RBTreeNode<E> parentOfu = (u == root) ? null :
182         (RBTreeNode<E>)(path.get(i - 1));
183     RBTreeNode<E> grandparentOfu = (parentOfu == null ||
184         parentOfu == root) ? null :
185         (RBTreeNode<E>)(path.get(i - 2));
186     RBTreeNode<E> childOfu = (u.left == null) ?
187         (RBTreeNode<E>)(u.right) : (RBTreeNode<E>)(u.left);
188
189     // Delete node u. Connect childOfu with parentOfu
190     connectNewParent(parentOfu, u, childOfu);
191
192     // Recolor the nodes and fix double black if needed
193     if (childOfu == root || u.isRed())
194         return; // Done if childOfu is root or if u is red
195     else if (childOfu != null && childOfu.isRed())
196         childOfu.setBlack(); // Set it black, done
197     else // u is black, childOfu is null or black
198         // Fix double black on parentOfu
199         fixDoubleBlack(grandparentOfu, parentOfu, childOfu, path, i);
200 }
201
202 /** Fix the double-black problem at node parent */
203 private void fixDoubleBlack(
204     RBTreeNode<E> grandparent, RBTreeNode<E> parent,
205     RBTreeNode<E> db, ArrayList<TreeNode<E>> path, int i) {
206     // Obtain y, y1, and y2
207     RBTreeNode<E> y = (parent.right == db) ?
208         (RBTreeNode<E>)(parent.left) : (RBTreeNode<E>)(parent.right);
209     RBTreeNode<E> y1 = (RBTreeNode<E>)(y.left);
210     RBTreeNode<E> y2 = (RBTreeNode<E>)(y.right);
211
212     if (y.isBlack() && y1 != null && y1.isRed()) {
213         if (parent.right == db) {
214             // Case 1.1: y is a left black sibling and y1 is red
215             connectNewParent(grandparent, parent, y);
216             recolor(parent, y, y1); // Adjust colors
217
218             // Adjust child links
219             parent.left = y.right;

```

```

220         y.right = parent;
221     }
222     else {
223         // Case 1.3: y is a right black sibling and y1 is red
224         connectNewParent(grandparent, parent, y1);
225         recolor(parent, y1, y); // Adjust colors
226
227         // Adjust child links
228         parent.right = y1.left;
229         y.left = y1.right;
230         y1.left = parent;
231         y1.right = y;
232     }
233 }
234 else if (y.isBlack() && y2 != null && y2.isRed()) {
235     if (parent.right == db) {
236         // Case 1.2: y is a left black sibling and y2 is red
237         connectNewParent(grandparent, parent, y2);
238         recolor(parent, y2, y); // Adjust colors
239
240         // Adjust child links
241         y.right = y2.left;
242         parent.left = y2.right;
243         y2.left = y;
244         y2.right = parent;
245     }
246     else {
247         // Case 1.4: y is a right black sibling and y2 is red
248         connectNewParent(grandparent, parent, y);
249         recolor(parent, y, y2); // Adjust colors
250
251         // Adjust child links
252         y.left = parent;
253         parent.right = y1;
254     }
255 }
256 else if (y.isBlack()) {
257     // Case 2: y is black and y's children are black or null
258     y.setRed(); // Change y to red
259     if (parent.isRed())
260         parent.setBlack(); // Done
261     else if (parent != root) {
262         // Propagate double black to the parent node
263         // Fix new appearance of double black recursively
264         db = parent;
265         parent = grandparent;
266         grandparent =
267             (i >= 3) ? (RBTreeNode<E>)(path.get(i - 3)) : null;
268         fixDoubleBlack(grandparent, parent, db, path, i - 1);
269     }
270 }
271 else { // y.isRed()
272     if (parent.right == db) {
273         // Case 3.1: y is a left red child of parent
274         parent.left = y2;
275         y.right = parent;
276     }
277     else {
278         // Case 3.2: y is a right red child of parent
279         parent.right = y.left;
280         y.left = parent;

```

```

281         }
282
283         parent.setRed(); // Color parent red
284         y.setBlack(); // Color y black
285         connectNewParent(grandparent, parent, y); // y is new parent
286         fixDoubleBlack(y, parent, db, path, i - 1);
287     }
288 }
289
290 /** Recolor parent, newParent, and c. Case 1 removal */
291 private void recolor(RBTreeNode<E> parent,
292     RBTreeNode<E> newParent, RBTreeNode<E> c) {
293     // Retain the parent's color for newParent
294     if (parent.isRed())
295         newParent.setRed();
296     else
297         newParent.setBlack();
298
299     // c and parent become the children of newParent; set them black
300     parent.setBlack();
301     c.setBlack();
302 }
303
304 /** Connect newParent with grandParent */
305 private void connectNewParent(RBTreeNode<E> grandparent,
306     RBTreeNode<E> parent, RBTreeNode<E> newParent) {
307     if (parent == root) {
308         root = newParent;
309         if (root != null)
310             newParent.setBlack();
311     }
312     else if (grandparent.left == parent)
313         grandparent.left = newParent;
314     else
315         grandparent.right = newParent;
316 }
317
318 @Override /** Preorder traversal from a subtree */
319 protected void preorder(TreeNode<E> root) {
320     if (root == null) return;
321     System.out.print(root.element +
322         ((RBTreeNode<E>)root).isRed() ? " (red) " : " (black) ");
323     preorder(root.left);
324     preorder(root.right);
325 }
326
327 /** RBTreeNode is TreeNode plus color indicator */
328 protected static class RBTreeNode<E> extends Comparable<E>> extends
329     BST.TreeNode<E> {
330     private boolean red = true; // Indicate node color
331
332     public RBTreeNode(E e) {
333         super(e);
334     }
335
336     public boolean isRed() {
337         return red;
338     }
339
340     public boolean isBlack() {
341         return !red;

```

```

342     }
343
344     public void setBlack() {
345         red = false;
346     }
347
348     public void setRed() {
349         red = true;
350     }
351
352     int blackHeight;
353 }
354 }

```

The **RBTree** class extends **BST**. Like the **BST** class, the **RBTree** class has a no-arg constructor that constructs an empty **RBTree** (lines 5–6) and a constructor that creates an initial **RBTree** from an array of elements (lines 9–11).

The **createNewNode()** method defined in the **BST** class creates a **TreeNode**. This method is overridden to return an **RBTreeNode** (lines 14–16). This method is invoked in the **insert** method in **BST** to create a node.

The **insert** method in **RBTree** is overridden in lines 20–29. The method first invokes the **insert** method in **BST**, then invokes **ensureRBTree(e)** (line 25) to ensure that tree is still a red-black tree after inserting a new element.

The **ensureRBTree(E e)** method first obtains the path of nodes that lead to element **e** from the root (line 34). It obtains **u** and **v** (the parent of **u**) from the path. If **u** is the root, color **u** black (lines 47–48). If **v** is red, invoke **fixDoubleRed** to fix the double red on both **u** and **v** (lines 49–50).

The **fixDoubleRed(u, v, path, i)** method fixes the double-red violation at node **u**. The method first obtains **w** (the grandparent of **u** from the path) (line 57), **parentOfw** if exists (lines 58–59), and **x** (the sibling of **v**) (lines 62–63). If **x** is **null** or black, consider four sub-cases to fix the double-red violation (lines 67–96). If **x** is red, color **w** and **u** red and color **w**'s two children black (lines 101–104). If **w** is the root, color **w** black (lines 104–106). Otherwise, propagate along the path to fix the new double-red violation (lines 109–111).

The **delete(E e)** method in **RBTree** is overridden in lines 134–174. The method locates the node that contains **e** (lines 136–146). If the node is null, no element is found (lines 148–149). The method considers two cases:

- If the node is internal, find the rightmost node in its left subtree (lines 156–159). Obtain a path from the root to the rightmost node (line 161), and replace the element in the node with the element in the rightmost node (line 164).
- If the node is external, obtain the path from the root to the node (line 167).

The last node in the path is the node to be deleted. Invoke **deleteLastNodeInPath(path)** to delete it and ensure the tree is a red-black after the node is deleted (line 170).

The **deleteLastNodeInPath(path)** method first obtains **u**, **parentOfu**, **grandparentOfu**, and **childOfu** (lines 180–187). **u** is the last node in the path. Connect **childOfu** as a child of **parentOfu** (line 190). This in effect deletes **u** from the tree. Consider three cases:

- If **childOfu** is the root or **childOfu** is red, we are done (lines 193–194).
- Otherwise, if **childOfu** is red, color it black (lines 195–196).
- Otherwise, invoke **fixDoubleBlack** to fix the double-black problem on **childOfu** (line 199).

The **fixDoubleBlack** method first obtains **y**, **y1**, and **y2** (lines 207–210). **y** is the sibling of the first double-black node, and **y1** and **y2** are the left and right children of **y**. Consider three cases:

- If **y** is black and **y1** or **y2** is red, fix the double-black problem for Case 1 (lines 213–255).
- Otherwise, if **y** is black, fix the double-black problem for Case 2 by recoloring the nodes. If parent is black and not a root, propagate double black to parent and recursively invoke **fixDoubleBlack** (lines 264–268).
- Otherwise, **y** is red. In this case, adjust the nodes to make parent the child of **y** (lines 272–281). Invoke **fixDoubleBlack** with the adjusted nodes (line 286) to fix the double-black problem.

The **preorder (TreeNode<E> root)** method is overridden to display the node colors (lines 319–325).

43.7 Testing the **RBTree** Class

*This section gives a test program that uses the **RBTree** class.*

Listing 43.4 gives a test program. The program creates an **RBTree** initialized with an array of integers **34**, **3**, and **50** (lines 4–5), inserts elements in lines 10–22, and deletes elements in lines 25–46.



LISTING 43.4 TestRBTree.java

```

1  public class TestRBTree {
2      public static void main(String[] args) {
3          // Create an RB tree
4          RBTree<Integer> tree =
5              new RBTree<Integer>(new Integer[] {34, 3, 50});
6          printTree(tree);
7
8          tree.insert(20);
9          printTree(tree);
10
11         tree.insert(15);
12         printTree(tree);
13
14         tree.insert(16);
15         printTree(tree);
16
17         tree.insert(25);
18         printTree(tree);
19
20         tree.insert(27);
21         printTree(tree);
22
23         tree.delete(50);
24         printTree(tree);
25
26         tree.delete(20);
27         printTree(tree);
28
29         tree.delete(15);
30         printTree(tree);
31     }

```

```

32     tree.delete(3);
33     printTree(tree);
34
35     tree.delete(25);
36     printTree(tree);
37
38     tree.delete(16);
39     printTree(tree);
40
41     tree.delete(34);
42     printTree(tree);
43
44     tree.delete(27);
45     printTree(tree);
46 }
47
48 public static <E extends Comparable<E>>
49     void printTree(BST <E> tree) {
50     // Traverse tree
51     System.out.print("\nInorder (sorted): ");
52     tree.inorder();
53     System.out.print("\nPostorder: ");
54     tree.postorder();
55     System.out.print("\nPreorder: ");
56     tree.preorder();
57     System.out.print("\nThe number of nodes is " + tree.getSize());
58     System.out.println();
59 }
60 }

```



```

Inorder (sorted): 3 34 50
Postorder: 3 50 34
Preorder: 34 (black) 3 (red) 50 (red)
The number of nodes is 3

Inorder (sorted): 3 20 34 50
Postorder: 20 3 50 34
Preorder: 34 (black) 3 (black) 20 (red) 50 (black)
The number of nodes is 4

Inorder (sorted): 3 15 20 34 50
Postorder: 3 20 15 50 34
Preorder: 34 (black) 15 (black) 3 (red) 20 (red) 50 (black)
The number of nodes is 5

Inorder (sorted): 3 15 16 20 34 50
Postorder: 3 16 20 15 50 34
Preorder: 34 (black) 15 (red) 3 (black) 20 (black) 16 (red) 50 (black)
The number of nodes is 6

Inorder (sorted): 3 15 16 20 25 34 50
Postorder: 3 16 25 20 15 50 34
Preorder: 34 (black) 15 (red) 3 (black) 20 (black) 16 (red) 25 (red)
50 (black)
The number of nodes is 7

Inorder (sorted): 3 15 16 20 25 27 34 50
Postorder: 3 16 15 27 25 50 34 20

```

```

Preorder: 20 (black) 15 (red) 3 (black) 16 (black) 34 (red) 25 (black)
          27 (red) 50 (black)
The number of nodes is 8

Inorder (sorted): 3 15 16 20 25 27 34
Postorder: 3 16 15 25 34 27 20
Preorder: 20 (black) 15 (red) 3 (black) 16 (black) 27 (red)
          25 (black) 34 (black)
The number of nodes is 7

Inorder (sorted): 3 15 16 25 27 34
Postorder: 3 15 25 34 27 16
Preorder: 16 (black) 15 (black) 3 (red) 27 (red) 25 (black) 34 (black)
The number of nodes is 6

Inorder (sorted): 3 16 25 27 34
Postorder: 3 25 34 27 16
Preorder: 16 (black) 3 (black) 27 (red) 25 (black) 34 (black)
The number of nodes is 5

Inorder (sorted): 16 25 27 34
Postorder: 25 16 34 27
Preorder: 27 (black) 16 (black) 25 (red) 34 (black)
The number of nodes is 4

Inorder (sorted): 16 27 34
Postorder: 16 34 27
Preorder: 27 (black) 16 (black) 34 (black)
The number of nodes is 3

Inorder (sorted): 27 34
Postorder: 34 27
Preorder: 27 (black) 34 (red)
The number of nodes is 2

Inorder (sorted): 27
Postorder: 27
Preorder: 27 (black)
The number of nodes is 1

Inorder (sorted):
Postorder:
Preorder:
The number of nodes is 0

```

Figure 43.14 shows how the tree evolves as elements are added to it, and Figure 43.28 shows how the tree evolves as elements are deleted from it.

43.8 Performance of the **RBTree** Class

This search, insertion, and deletion operations take $O(\log n)$ time in a red-black tree.

The search, insertion, and deletion times in a red-black tree depend on the height of the tree. A red-black tree corresponds to a 2–4 tree. When you convert a node in a 2–4 tree to red-black tree nodes, you get one black node and zero, one, or two red nodes as its children, depending on whether the original node is a 2-node, 3-node, or 4-node. So, the height of a red-black tree



TABLE 43.1 Time Complexities for Methods in **RBTree**, **AVLTree**, and **Tree234**

Methods	Red-Black Tree	AVL Tree	2-4 Tree
search (e: E)	$O(\log n)$	$O(\log n)$	$O(\log n)$
insert (e: E)	$O(\log n)$	$O(\log n)$	$O(\log n)$
delete (e: E)	$O(\log n)$	$O(\log n)$	$O(\log n)$
getSize()	$O(1)$	$O(1)$	$O(1)$
isEmpty()	$O(1)$	$O(1)$	$O(1)$

is at most as twice that of its corresponding 2–4 tree. Since the height of a 2–4 tree is $\log n$, the height of a red-black tree is $2\log n$.

A red-black tree has the same time complexity as an AVL tree, as shown in Table 43.1. In general, a red-black is more efficient than an AVL tree, because a red-black tree requires only one-time restructuring of the nodes for insert and delete operations.

A red-black tree has the same time complexity as a 2–4 tree, as shown in Table 43.1. In general, a red-black is more efficient than a 2–4 tree for two reasons:

1. A red-black tree requires only one-time restructuring of the nodes for insert and delete operations. However, a 2–4 tree may require many splits for an insert operation and fusion for a delete operation.
2. A red-black tree is a binary search tree. A binary tree can be implemented more space efficiently than a 2–4 tree, because a node in a 2–4 tree has at most three elements and four children. Space is wasted for 2-nodes and 3-nodes in a 2–4 tree.

Listing 43.5 gives an empirical test of the performance of AVL trees, 2–4 trees, and red-black trees.

LISTING 43.5 TreePerformanceTest.java

```
1 public class TreePerformanceTest {
2     public static void main(String[] args) {
3         final int TEST_SIZE = 500000; // Tree size used in the test
4
5         // Create an AVL tree
6         Tree<Integer> tree1 = new AVLTree<Integer>();
7         System.out.println("AVL tree time: " +
8             getTime(tree1, TEST_SIZE) + " milliseconds");
9
10        // Create a 2-4 tree
11        Tree<Integer> tree2 = new Tree24<Integer>();
12        System.out.println("2-4 tree time: "
13            + getTime(tree2, TEST_SIZE) + " milliseconds");
14
15        // Create a red-black tree
16        Tree<Integer> tree3 = new RBTree<Integer>();
17        System.out.println("RB tree time: "
18            + getTime(tree3, TEST_SIZE) + " milliseconds");
19    }
20
21    public static long getTime(Tree<Integer> tree, int testSize) {
22        long startTime = System.currentTimeMillis(); // Start time
23
24        // Create a list to store distinct integers
25        java.util.List<Integer> list = new java.util.ArrayList<Integer>();
```

```

26     for (int i = 0; i < testSize; i++)
27         list.add(i);
28
29     java.util.Collections.shuffle(list); // Shuffle the list
30
31     // Insert elements in the list to the tree
32     for (int i = 0; i < testSize; i++)
33         tree.insert(list.get(i));
34
35     java.util.Collections.shuffle(list); // Shuffle the list
36
37     // Delete elements in the list from the tree
38     for (int i = 0; i < testSize; i++)
39         tree.delete(list.get(i));
40
41     // Return elapse time
42     return System.currentTimeMillis() - startTime;
43 }
44 }

```

```

AVL tree time: 7609 milliseconds
2-4 tree time: 8594 milliseconds
RB tree time: 5515 milliseconds

```



The **getTestTime** method creates a list of distinct integers from **0** to **testSize - 1** (lines 25–27), shuffles the list (line 29), adds the elements from the list to a tree (lines 32–33), shuffles the list again (line 35), removes the elements from the tree (lines 38–39), and finally returns the execution time (line 42).

The program creates an AVL (line 6), a 2-4 tree (line 11), and a red-black tree (line 16). The program obtains the execution time for adding and removing **500000** elements in the three trees.

As you see, the red-black tree performs the best, followed by the AVL tree.



Note

The **java.util.TreeSet** class in the Java API is implemented using a red-black tree. Each entry in the set is stored in the tree. Since the **search**, **insert**, and **delete** methods in a red-black tree take $O(\log n)$ time, the **get**, **add**, **remove**, and **contains** methods in **java.util.TreeSet** take $O(\log n)$ time.



Note

The **java.util.TreeMap** class in the Java API is implemented using a red-black tree. Each entry in the map is stored in the tree. The order of the entries is determined by their keys. Since the **search**, **insert**, and **delete** methods in a red-black tree take $O(\log n)$ time, the **get**, **put**, **remove**, and **containsKey** methods in **java.util.TreeMap** take $O(\log n)$ time.

KEY TERMS

black depth	43-2	external node	43-9
double-black violation	43-11	red-black tree	43-2
double-red violation	43-7		

CHAPTER SUMMARY

1. A red-black tree is a binary search tree, derived from a *2-4 tree*. A red-black tree corresponds to a 2-4 tree. You can convert a red-black tree to a 2-4 tree or vice versa.
2. In a red-black tree, each node is colored red or black. The root is always black. Two adjacent nodes cannot be both red. All external nodes have the same black depth.
3. Since a red-black tree is a binary search tree, the `RBTree` class extends the `BST` class.
4. Searching an element in a red-black tree is the same as in binary search tree, since a red-black tree is a binary search tree.
5. A new element is always inserted as a leaf node. If the new node is the root, color it black. Otherwise, color it red. If the parent of the new node is red, we have to fix the *double-red violation* by reassigning the color and/or restructuring the tree.
6. If a node to be deleted is internal, find the rightmost node in its left subtree. Replace the element in the node with the element in the rightmost node. Delete the rightmost node.
7. If the external node to be deleted is red, simply reconnect the parent node of the external node with the child node of the external node.
8. If the external node to be deleted is black, you need to consider several cases to ensure that black height for external nodes in the tree is maintained correctly.
9. The height of a red-black tree is $O(\log n)$. So, the time complexities for the `search`, `insert`, and `delete` methods are $O(\log n)$.



Quiz

Answer the quiz for this chapter online at the book Companion Website.

MyProgrammingLab™

PROGRAMMING EXERCISES

- *43.1** (*red-black tree to 2-4 tree*) Write a program that converts a red-black tree to a 2-4 tree.
- *43.2** (*2-4 tree to red-black tree*) Write a program that converts a red-black tree to a 2-4 tree.
- ***43.3** (*red-black tree animation*) Write a GUI program that animates the red-black tree `insert`, `delete`, and `search` methods, as shown in Figure 43.6.
- **43.4** (*Parent reference for RBTree*) Suppose the `TreeNode` class defined in `BST` contains a reference to the node's parent, as shown in Exercise 26.17. Implement the `RBTree` class to support this change. Write a test program that adds numbers 1, 2, . . . , 100 to the tree and displays the paths for all leaf nodes.